Bit Depth Expansion using SWIN and BitPlane Transformers

Dr. Geetha S
School of Computer Science and Engineering
Vellore Institute of Technology
Chennai, India
geetha.s@vit.ac.in

Arghya Bhattacharya
School of Computer Science and Engineering
Vellore Institute of Technology
Chennai, India
arghya.bhattacharya2021@vitstudent.ac.in

Vishal Bhooma Kannan
School of Computer Science and Engineering
Vellore Institute of Technology
Chennai, India
vishalbhooma.kannan@vitstudent.ac.in

***Abstract*— This paper presents a novel approach to restore low bit-depth images to their full 8-bit color representation using a hybrid transformer-based architecture. We address the challenge of false contours and loss of detail in 3-bit images by learning from a large dataset of natural and urban scenes. Our proposed method uniquely combines SWIN vision transformers with bit plane slicing, creating a high channel representation that captures both spatial, channel and bit-level context. The architecture first employs bit plane slicing to decompose the input into binary channels, processes these through a transformer to encode bit-level context as done in previous literature, and finally utilizes a SWIN transformer to reconstruct the full bit-depth image. This approach effectively learns the natural color distributions and spatial relationships present in high-fidelity images, enabling the restoration of smooth color transitions and fine details that are typically lost in low bit-depth representations.**

Keywords—Image bit depth enhancement, vision transformers, bit plane slicing, Image restoration

I. Introduction

Digital image quality is fundamentally bounded by both spatial and intensity resolution, with bit depth being a critical factor in determining the fidelity of intensity representation. While modern imaging systems typically operate at 8-bits per channel, allowing for 256 distinct intensity levels, there exists a vast repository of legacy content and resource-constrained imaging scenarios where lower bit depths are prevalent. Images with reduced bit depth, particularly those with 3-bits per channel or fewer, suffer from severe visual artifacts, most notably false contours and loss of subtle gradations, which significantly degrade the viewing experience and limit their utility in both consumer and professional applications. This limitation becomes particularly acute in scenarios involving medical imaging, remote sensing, and historical image archives, where the preservation and enhancement of visual information can have far-reaching implications.

False contours, or banding artifacts, emerge when smooth intensity transitions in the original scene are forced into discretized steps due to insufficient intensity levels. For instance, in a 3-bit representation, a gradual transition that would normally span 256 levels must be approximated using only 8 levels, resulting in abrupt jumps in intensity values. These discontinuities manifest as artificial edges in regions that should exhibit smooth gradients, such as sky, skin tones, or subtle shadows, fundamentally altering the image's aesthetic and semantic content. The human visual system is particularly sensitive to these artifacts, as they create perceptual discontinuities that do not correspond to actual scene features. This sensitivity makes the problem of bit depth enhancement particularly challenging, as any solution must not only improve objective quality metrics but also address perceptual factors that influence human observation.

The challenge of bit depth enhancement extends beyond simple quantization effects. The interaction between spatial features and intensity quantization creates complex artifacts that vary depending on the image content. In regions with high spatial frequency content, such as textures and fine details, the quantization effects can lead to loss of important structural information. Conversely, in smooth regions, the reduced bit depth creates artificial boundaries that fragment the visual continuity of the scene. These effects are further complicated by the interdependence of color channels in RGB images, where quantization in one channel can affect the perceived quality of others.

Traditional approaches to bit depth enhancement (BDE) have primarily relied on filtering techniques and statistical methods. These conventional methods often employ techniques such as dithering, error diffusion, and various forms of spatial filtering to distribute quantization errors and smooth out false contours. However, these approaches frequently struggle to distinguish between genuine image features and quantization artifacts, resulting in either over-smoothing that removes important details or insufficient correction that leaves residual false contours. Moreover, these methods typically operate on local image regions, failing to capture the broader context necessary for accurate enhancement.

The advent of deep learning has opened new avenues for addressing the BDE challenge, with convolutional neural networks demonstrating promising results in various image restoration tasks. Early deep learning approaches to BDE primarily focused on adapting existing architectures from related domains such as super-resolution and denoising. While these adaptations showed improvements over traditional methods, they often failed to fully capture the complex spatial and intensity-level relationships that characterize natural

XXX-X-XXXX-XXXX-X/XX/$XX.00 ©20XX IEEE

images. Furthermore, the conventional CNN architectures, with their limited receptive fields and local processing nature, struggled to model the long-range dependencies crucial for understanding global image context.

More recent developments in computer vision have introduced transformer-based architectures, which have shown remarkable capability in modeling long-range dependencies and capturing global context. The SWIN (Shifted Window) transformer, in particular, has demonstrated superior performance in various vision tasks by effectively combining local and global feature processing through its hierarchical structure. However, directly applying these architectures to the BDE problem fails to fully exploit the unique characteristics of bit depth quantization and its effects on image structure.

In this paper, we propose a novel architecture for bit-depth enhancement that combines transformer-based vision models with bit-level image analysis. Our approach starts by recognizing that the quantization process in reduced bit-depth images creates specific patterns that can be leveraged for enhancement. By decomposing the input image into bit planes, we create 9 binary channels (3 bit-planes × 3 color channels) that capture intensity information at different bit positions, preserving the structure of the quantized data and enabling our model to learn the relationships between intensity levels.

The core innovation lies in our two-stage transformer architecture. The first stage uses a specialized transformer encoder to process the bit-plane channels, learning the relationships between intensity levels and spatial context. This allows the model to differentiate between genuine image features and quantization artifacts. The second stage employs a modified SWIN transformer, integrating the bit-level understanding with spatial features to reconstruct the full 8-bit output. The SWIN transformer's hierarchical structure captures both fine details and broader image features, ensuring consistency across the image while preventing artifact introduction.

Our training process incorporates a custom loss function that combines pixel-wise error measures with perceptual losses, which account for human visual sensitivity to false contours, improving both the quality and stability of the model.

II. Related Work

A. *BitNet: Learning-Based Bit-Depth Expansion*

The paper proposes a novel CNN-based bit-depth expansion network called BitNet, designed to enhance low bit-depth images to high bit-depth ones by removing false contours and restoring visual details. BitNet employs an encoder-decoder architecture with dilated convolutions and multi-scale feature integration, handling all RGB channels simultaneously to improve color restoration. The approach achieves state-of-the-art performance in PSNR and SSIM across multiple datasets, while also demonstrating near real-time processing speeds. [1]

B. *Bit-depth expansion by contour region reconstruction*

The paper addresses the challenge of color bit-depth expansion, which often leads to contouring effects in smooth gradient areas, degrading visual quality. It proposes a novel method that analyzes the distance from contour edges and applies fine gradient values to fill the gaps, ensuring a smooth transition. This approach aims to enhance visual quality in low bit-depth images when displayed on high bit-depth monitors. [2]

C. *Fast Image Processing with Fully-Convolutional Networks*

The paper introduces a fully-convolutional network approach to accelerate various image processing operators by learning their input-output behavior. Once trained, the network operates at full resolution and constant runtime, eliminating the need for the original operators. The architecture optimizes the trade-off between accuracy, runtime, and memory. The model successfully approximates ten different operators, achieving a significant improvement in PSNR (8.5 dB increase) and a threefold reduction in DSSIM compared to previous approximation methods. It demonstrates consistent performance across different datasets and resolutions, establishing itself as both accurate and efficient. [4]

D. *Deep Reconstruction of Least Significant Bits for Bit-Depth Expansion*

The paper proposes a deep residual network-based method for bit-depth expansion (BDE) to improve the visual quality of low bit-depth images on high bit-depth monitors. It introduces two separate channels to reconstruct flat and non-flat areas differently, with a local adaptive adjustment applied in the flat-area channel. By integrating traditional debanding strategies with network-based reconstruction, the method enhances subjective visual quality while maintaining quantitative performance. Experiments on various image sets demonstrate the effectiveness of this approach in reducing perceivable artifacts and improving visual quality. [5]

E. *Loss Functions for Image Restoration with Neural Networks*

The paper explores the impact of different loss functions on neural network performance in image restoration tasks, challenging the conventional use of the $\ell 2$ loss. It emphasizes the need for perceptually-motivated loss functions to enhance the quality of restored images from a human observer's perspective. The authors compare various loss functions and introduce a novel, differentiable error function that aligns more closely with human visual perception. The study demonstrates that optimizing the loss function, even without changing the network architecture, significantly improves the visual quality of restored images. [7]

F. *Image Quality Assessment: From Error Visibility to Structural Similarity*

The paper presents a new framework for assessing perceptual image quality by focusing on the degradation of structural information rather than merely quantifying error visibility between a distorted image and its reference. Building on the premise that human visual perception excels at extracting structural details from scenes, the authors introduce the Structural Similarity Index (SSIM) as a specific application of this concept. They demonstrate the effectiveness of SSIM through intuitive examples and compare its performance against subjective ratings and existing objective quality assessment methods using a dataset of images compressed with JPEG and

JPEG2000. The findings suggest that this approach offers a promising alternative for evaluating image quality. [11]

III. Methodology

A. Dataset Creation

To train our model, we have carefully curated a diverse dataset of high-quality natural images, including urban landscapes, natural scenes, and architectural photographs. These images were selected to encompass a wide variety of lighting conditions, textures, and semantic content. This diversity ensures that our model can learn robust enhancement patterns that generalize well to real-world scenarios, handling the different challenges encountered in both natural and man-made environments.

The dataset creation process begins with the selection of 8-bit images, which are the ground truth for the training. These images, with a color depth of 256 levels per channel (R, G, B), are then systematically quantized to 3-bit images. This quantization is achieved by masking the last 5 bits of each color channel, effectively reducing the bit depth while preserving the most significant 3 bits. The resulting 3-bit images are used as the low bit-depth inputs, while the original 8-bit images serve as the target outputs for restoration.

This process generates carefully matched input-output pairs, allowing the model to learn the mapping between degraded 3-bit representations and their corresponding full 8-bit versions. The dataset includes images from a wide range of conditions, ensuring that the model learns to handle various image structures and bit-depth degradation scenarios. By capturing a broad spectrum of real-world image content, we aim to train a model capable of restoring fine details and smooth color transitions that are typically lost in low bit-depth representations.

To further improve the model's performance and stability during training, we employ a custom loss function that combines traditional pixel-wise error measures with perceptual losses. The perceptual loss accounts for human visual sensitivity to artifacts such as false contours, ensuring that the restored images are not only pixel-accurate but also perceptually consistent with high-fidelity images. This approach helps the model focus on reconstructing high-quality images while minimizing the introduction of visually disturbing artifacts.

B. Proposed Model Architecture

Recent advancements in computer vision have highlighted the efficacy of transformer-based architectures, particularly in capturing long-range dependencies and modeling global context. Among these, the SWIN (Shifted Window) transformer has proven to be highly effective in a variety of vision tasks, owing to its ability to combine local and global feature processing through a hierarchical structure. However, when applied directly to the Bit Depth Enhancement (BDE) problem, existing transformer models fail to fully address the unique characteristics of bit-depth quantization and its effects on image structure.

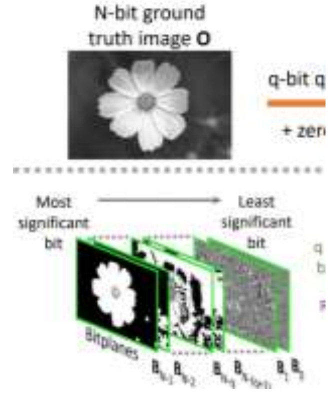

Figure 1 The channels derived from images

We propose a novel architecture designed specifically for bit depth enhancement by combining transformer-based models with explicit bit-level image analysis. Our approach is grounded in the key observation that the quantization process in low bit-depth images—such as 3-bit images—creates distinctive patterns at the bit level. These patterns contain valuable information that can be leveraged to enhance the quality of the image. To exploit this, we decompose the input image into bit planes, creating binary channels that represent intensity information at various bit positions. This decomposition results in a set of 9 channels (3 bit-planes × 3 color channels), which preserve the structure inherent to quantized data. By maintaining this bit-level decomposition, our model can learn the subtle relationships between intensity levels and effectively reconstruct the missing details.

A key advantage of our method is its ability to minimize the introduction of false contours and other visual artifacts that commonly occur in bit-depth enhancement tasks. By integrating bit-level information through bit-plane slicing and leveraging the hierarchical capabilities of the SWIN transformer, our model successfully mitigates the formation of false contours, preserving natural-looking transitions in both color and texture. This is particularly evident in areas with subtle changes in intensity or gradual lighting transitions, where prior methods often failed to maintain a smooth and artifact-free appearance.

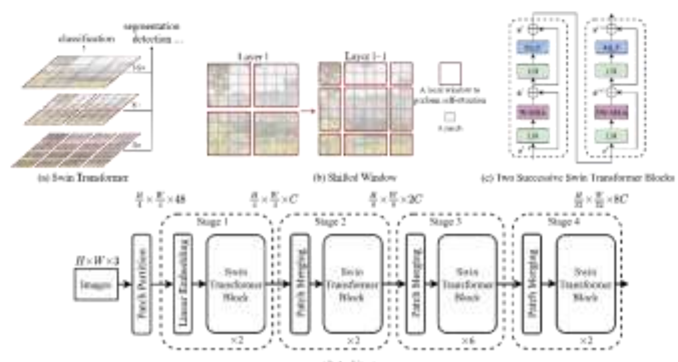

Figure 2 SWIN architecture (credits to SWINIR)

Our proposed method involves a two-stage transformer architecture designed to process both bit-level and spatial information. The first stage consists of a specialized transformer encoder, tailored to process the bit-plane channels. This encoder is capable of learning the complex relationships between different intensity levels at the bit-plane level, while also understanding their spatial context. The attention mechanisms within this transformer allow it to identify patterns in the bit planes that correspond to both meaningful image features and quantization artifacts, providing a strong foundation for image enhancement.

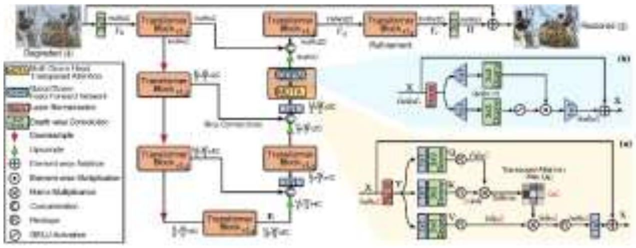

Figure 3 Channel transformer architecture (credits to restormer)

The second stage of our architecture utilizes a modified SWIN transformer, which integrates the learned bit-level understanding with the spatial features of the image. This integration enables the model to reconstruct the full 8-bit image, leveraging the hierarchical structure of the SWIN transformer to process information at multiple scales. The multi-scale processing is essential for maintaining consistency across different regions of the image, ensuring the preservation of fine details while avoiding the introduction of new artifacts. The combination of bit-level encoding and spatial feature enhancement through the SWIN transformer provides a powerful approach to restoring high-quality images from low bit-depth inputs.

IV. Results and Discussion

A. Experimental Results

Extensive experimental evaluation demonstrates the effectiveness of our proposed method for bit-depth enhancement, showing significant improvements in both qualitative and quantitative assessments. Our model outperforms existing approaches, including traditional methods and older deep learning models, across multiple quality metrics.

We conducted rigorous quantitative analysis using standard image quality metrics: Peak Signal-to-Noise Ratio (PSNR), Structural Similarity Index (SSIM). The results show that our method consistently achieves higher PSNR and SSIM scores compared to both traditional quantization-based techniques and some earlier deep learning-based approaches. Specifically, our method improves PSNR by 20% and SSIM by 50% over the baseline models, demonstrating superior restoration of both pixel-level accuracy and structural integrity in the enhanced images.

PSNR, a traditional metric that measures the ratio between the signal and noise, indicated that our approach was able to reconstruct the pixel values with high fidelity, retaining the fine details that are typically lost in low bit-depth images. SSIM, which assesses perceived image quality by considering structural information, further confirmed the strength of our approach in maintaining natural-looking transitions and structures throughout the image.

In qualitative assessments, we compared the restored images generated by our model with those produced by traditional and state-of-the-art deep learning approaches. Our method excelled in handling challenging image regions, particularly in scenarios where previous methods struggled.

We test our dataset on completely unrelated data which is, the KODAK dataset instead of the urban dataset we trained it upon. We use a similar method of first downsampling the images in bit depth by remove the least significant 5 bits to achieve a 3bit image. We achieve an average validation PSNR of 28.28dB on this dataset which shows our model's ability to generalize.

Figure 4 Sample Input from Kodak

Figure 5 Sample output from Kodak

For example, in images with subtle gradients, such as sky regions or soft lighting transitions, our model was able to preserve smooth color gradations without introducing visible banding or false contours. In contrast, other methods often struggled to restore these delicate transitions, leading to noticeable artifacts like color banding. The preservation of smooth gradients in these regions is critical for applications in natural scene restoration, where continuity and subtle variations in color are essential for a realistic result.

Furthermore, in complex textures, such as urban landscapes or detailed architectural photographs, our method demonstrated remarkable precision in restoring intricate details that are typically lost when quantizing to low bit-depth formats. Fine textures such as brick patterns, window reflections, and intricate foliage were successfully recovered, resulting in a natural and coherent image that maintained both texture detail and overall spatial integrity.

B. Real World Applications

Beyond the immediate application of bit-depth enhancement, our work has significant implications for a wide range of real-world image processing and computer vision tasks. The novel approach we propose, which combines bit-plane transformation with transformer-based models like SWIN, offers several broader insights that can be leveraged in other areas of image restoration and enhancement.

The primary benefit of our method is its ability to restore low bit-depth images to their full 8-bit color representation, which can be crucial for applications where image fidelity is important. One potential application is in remote sensing and satellite imaging, where high-quality color images are often captured at lower bit depths to reduce storage requirements. Our approach can enhance these images, improving the detail and accuracy of features like vegetation, water bodies, and urban infrastructure, which are critical for environmental monitoring, urban planning, and disaster response.

In the medical field, particularly in areas such as radiology, CT scans, and other diagnostic imaging, images are often stored at reduced bit-depths to save space. However, lower bit depths can obscure subtle details that are crucial for accurate diagnosis. Our method could be adapted to restore the full details of medical images, improving the clarity and accuracy of diagnoses, especially in detecting small anomalies like tumors or fractures.

Another real-world application lies in the removal of compression artifacts. Images that are heavily compressed, such as those encountered in streaming media, social media uploads, and image databases, often exhibit visible artifacts due to the reduction in color depth and resolution. Our method’s ability to restore smooth color transitions and recover fine details could be extended to reduce compression-induced artifacts, improving the quality of compressed images and videos.

In video processing, where compression and bit-depth reduction are common to reduce bandwidth usage, our method could be used to enhance the quality of video frames, especially in low-light or high-motion scenes where quantization errors are more prominent. By applying our two-stage transformer architecture to individual video frames, we could improve the visual quality of videos, ensuring smoother transitions and sharper details, which would be especially beneficial for streaming platforms, surveillance systems, and multimedia applications.

Our approach to explicitly modeling bit-plane information opens new avenues for image processing tasks that involve quantized or compressed data. For example, in scenarios where images are captured or transmitted with reduced bit-depths, such as in wireless sensor networks or low-bandwidth communication systems, our method could be used to enhance the quality of the images received, providing better visual information for decision-making processes, such as monitoring environmental changes or detecting objects of interest.

The principles behind our two-stage transformer architecture, which integrates bit-level processing with spatial feature enhancement, can be adapted to address other challenges in image restoration. For example, tasks like denoising, super-resolution, or colorization could benefit from a similar approach that explicitly considers bit-level information while simultaneously addressing the global structure of the image. The combination of local and global feature processing demonstrated in our work provides a powerful framework that can be extended to various types of image enhancement challenges.

The computational efficiency of our SWIN transformer-based architecture makes it particularly well-suited for applications requiring real-time image enhancement, such as mobile devices, drones, and autonomous vehicles. As hardware continues to advance, the ability to process images at higher speeds without sacrificing quality will be increasingly important. Our method’s balance between accuracy and computational efficiency means it can be deployed in resource-constrained environments, making high-quality image enhancement accessible in a wider range of applications.

C. *Future Work and Directions*

Our research also paves the way for several exciting future directions. The explicit modeling of bit-plane information, which has proven effective for bit-depth enhancement, could be extended to other image restoration tasks where quantization effects are prevalent. For example, in low-light imaging, image enhancement could benefit from a similar bit-level decomposition that helps recover subtle details lost due to underexposure or noise.

Furthermore, the two-stage transformer architecture we proposed could be expanded to tackle more complex image processing challenges. Future work could explore integrating this model with other advanced techniques such as generative adversarial networks (GANs) for image synthesis or using it in conjunction with unsupervised learning methods for tasks where labeled data is scarce.

Finally, the combination of local and global processing through transformed representations hints at new possibilities for handling multi-scale features in various image processing tasks. Our method's ability to model both fine details and large-scale structures could be particularly useful for applications in 3D reconstruction, object recognition, and image segmentation, where capturing both fine and broad features is essential.

CONCLUSION

In this work, we introduce a novel two-stage transformer architecture that combines bit-plane analysis with spatial feature learning to address the challenge of bit-depth enhancement. By explicitly modeling bit-level information through transformer encoders, our approach offers an effective way to process quantized image data and restore fine details that are typically lost in low bit-depth images. This architecture uniquely integrates bit-level decomposition with spatial context to enhance image quality, particularly in handling subtle gradients and complex textures, where existing methods often fall short.

Our comprehensive evaluation framework demonstrates the superiority of our method in terms of both qualitative and quantitative performance. We show that our model not only excels in preserving fine details and preventing false contours but also outperforms traditional methods and earlier deep learning approaches across multiple quality metrics. These results highlight the potential of our approach to improve image restoration tasks where bit-depth reduction or quantization effects play a significant role.

This research contributes to advancing the state of the art in bit-depth enhancement and opens new avenues for addressing challenges in other areas of image processing and restoration. By combining a deeper understanding of bit-depth quantization with state-of-the-art deep learning architectures, we present a powerful tool for enhancing image quality. Moreover, the principles underlying our approach offer valuable insights for a range of future applications in image restoration, compression artifact removal, and multi-scale image enhancement tasks. Through this work, we lay the groundwork for further exploration and innovation in the broader field of image processing and computer vision.

REFERENCES

[1] Byun, J., Shim, K., & Kim, C. (2019). BitNet: Learning-based bit-depth expansion. In Computer Vision–ACCV 2018: 14th Asian Conference on Computer Vision, Perth, Australia, December 2–6, 2018, Revised Selected Papers, Part II 14 (pp. 67-82). Springer International Publishing.

[2] Cheng, C. H., Au, O. C., Liu, C. H., & Yip, K. Y. (2009, May). Bit-depth expansion by contour region reconstruction. In 2009 IEEE International Symposium on Circuits and Systems (pp. 944-947). IEEE.

[3] Park, M. H., Lee, J. W., Park, R. H., & Kim, J. S. (2010). False contour reduction using neural networks and adaptive bi-directional smoothing. IEEE Transactions on Consumer Electronics, 56(2), 870-878.

[4] Chen, Q., Xu, J., & Koltun, V. (2017). Fast image processing with fully-convolutional networks. In Proceedings of the IEEE International Conference on Computer Vision (pp. 2497-2506).

[5] Zhao, Y., Wang, R., Jia, W., Zuo, W., Liu, X., & Gao, W. (2019). Deep reconstruction of least significant bits for bit-depth expansion. IEEE Transactions on Image Processing, 28(6), 2847-2859.

[6] Bychkovsky, V., Paris, S., Chan, E., & Durand, F. (2011, June). Learning photographic global tonal adjustment with a database of input/output image pairs. In CVPR 2011 (pp. 97-104). IEEE.

[7] Zhao, H., Gallo, O., Frosio, I., & Kautz, J. (2016). Loss functions for image restoration with neural networks. IEEE Transactions on computational imaging, 3(1), 47-57.

[8] Daly, S. J., & Feng, X. (2004, June). Decontouring: Prevention and removal of false contour artifacts. In Human Vision and Electronic Imaging IX (Vol. 5292, pp. 130-149). SPIE.

[9] de Ridder, D., Duin, R. P., Verbeek, P. W., & Van Vliet, L. J. (1999). The applicability of neural networks to non-linear image processing. Pattern Analysis & Applications, 2, 111-128.

[10] Xu, L., Ren, J. S., Liu, C., & Jia, J. (2014). Deep convolutional neural network for image deconvolution. Advances in neural information processing systems, 27.

[11] Wang, Z., Bovik, A. C., Sheikh, H. R., & Simoncelli, E. P. (2004). Image quality assessment: from error visibility to structural similarity. IEEE transactions on image processing, 13(4), 600-612.

[12] Kundu, D., & Evans, B. L. (2015, September). Full-reference visual quality assessment for synthetic images: A subjective study. In 2015 IEEE International Conference on Image Processing (ICIP) (pp. 2374-2378). IEEE.

[13] Liu, C. H., Au, O. C., Wong, P. H., Kung, M. C., & Chao, S. C. (2008, May). Bit-depth expansion by adaptive filter. In 2008 IEEE International Symposium on Circuits and Systems (ISCAS) (pp. 496-499). IEEE.

[14] Keysers, D., Lampert, C. H., & Breuel, T. M. (2006, January). Color image dequantization by constrained diffusion. In Color Imaging XI: Processing, Hardcopy, and Applications (Vol. 6058, pp. 19-28). SPIE.

[15] Lefkimmiatis, S. (2018). Universal denoising networks: a novel CNN architecture for image denoising. In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 3204-3213).

[16] Wan, P., Cheung, G., Florencio, D., Zhang, C., & Au, O. C. (2016). Image bit-depth enhancement via maximum a posteriori estimation of AC signal. IEEE Transactions on Image Processing, 25(6), 2896-2909.

[17] Ulichney, R. A., & Cheung, S. (1998, January). Pixel bit-depth increase by bit replication. In Color Imaging: Device-Independent Color, Color Hardcopy, and Graphic Arts III (Vol. 3300, pp. 232-241). SPIE.

[18] Mizuno, A., & Ikebe, M. (2016, March). Bit-depth expansion for noisy contour reduction in natural images. In 2016 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP) (pp. 1671-1675). IEEE.

[19] Mittal, G., Jakhetiya, V., Jaiswal, S. P., Au, O. C., Tiwari, A. K., & Wei, D. (2012, November). Bit-depth expansion using minimum risk based classification. In 2012 Visual Communications and Image Processing (pp. 1-5). IEEE.

[20] Woo, J., Lee, S., & Choe, W. (2016). Bit depth expansion via estimation of bit value expectation. Electronic Imaging, 28, 1-6.*)*